


MACHINE LEARNING

COMP8220

02 – Workflow for ML Project



❖ Workflow of Machine Learning Project

- Data Pre-processing and Feature Engineering
- Model Training and Evaluation

❖ Linear Regression Introduction

- Linear Regression Model
- Overfitting and Model Selection

- ❖ Workflow of Machine Learning Project
 - Data Pre-processing and Feature Engineering
 - Model Training and Evaluation

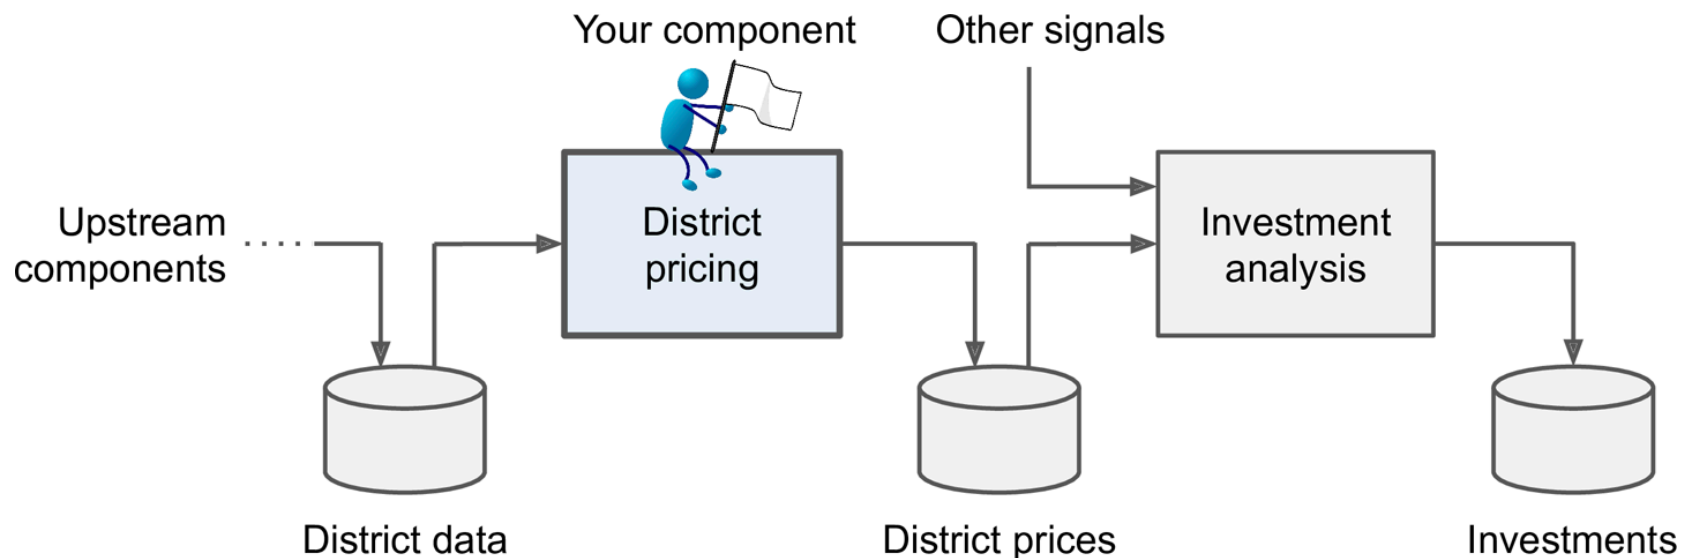
Workflow of ML Project



MACQUARIE
University


- ❖ Problem framing
- ❖ Data collection
- ❖ Exploratory analysis and visualization
- ❖ Data Pre-processing
- ❖ Feature extraction and selection
- ❖ ML model training
- ❖ Model fine-tuning
- ❖ ML model deployment and maintenance

- ❖ What is the business objective?
 - Building a model might not be the end goal
 - E.g., an ML pipeline for real estate investment



- ❖ Requirement analysis
 - Similar to that in Software Engineering
 - Based on data and business requirement
- ❖ Typical examples of analysis questions
 - Is it a supervised or unsupervised learning task?
 - Are data labelled or not?
 - Will domain experts be involved in any stage?
 - Is it a online learning or batch learning?
 - Data streaming? Large datasets?
 - What is the performance measures?
 - Sequential or parallel implementation?
 - What are the interfaces to other components?

- ❖ Manually collect datasets (usually small)
- ❖ Automated tools to collect datasets
 - E.g., Web crawlers to collect data from web pages
- ❖ Public datasets
 - Data repositories, e.g.,
 - UC Irvine Machine Learning Repository
 - Kaggle datasets
 - Amazon's AWS datasets
 - Data Portals (meta portals listing repositories)
 - Web APIs for data queries, e.g.,
 - Twitter data APIs
 - Domain property data APIs



- ❖ Quick glance to get a general understanding
 - Metadata, e.g., data size, data types, etc.
 - Assist in problem framing and modelling
 - Less complex than machine learning models
 - Light-weight experiments (feature combination)
- ❖ Visualization
 - Human brains are very good at spotting patterns in pictures
 - Applicable to low (2 or 3) dimensional data
 - Dimension reduction might be necessary for visualization
 - PCA
 - Just randomly select 2 or 3 features for visualization (there are many possible combinations)

❖ Typical types of exploratory analysis



- Statistics
- Correlations

❖ Typical statistics

- Count
- Mean, median, mode
- Std
- Max, min
- Quantile
- ...

❖ Histogram

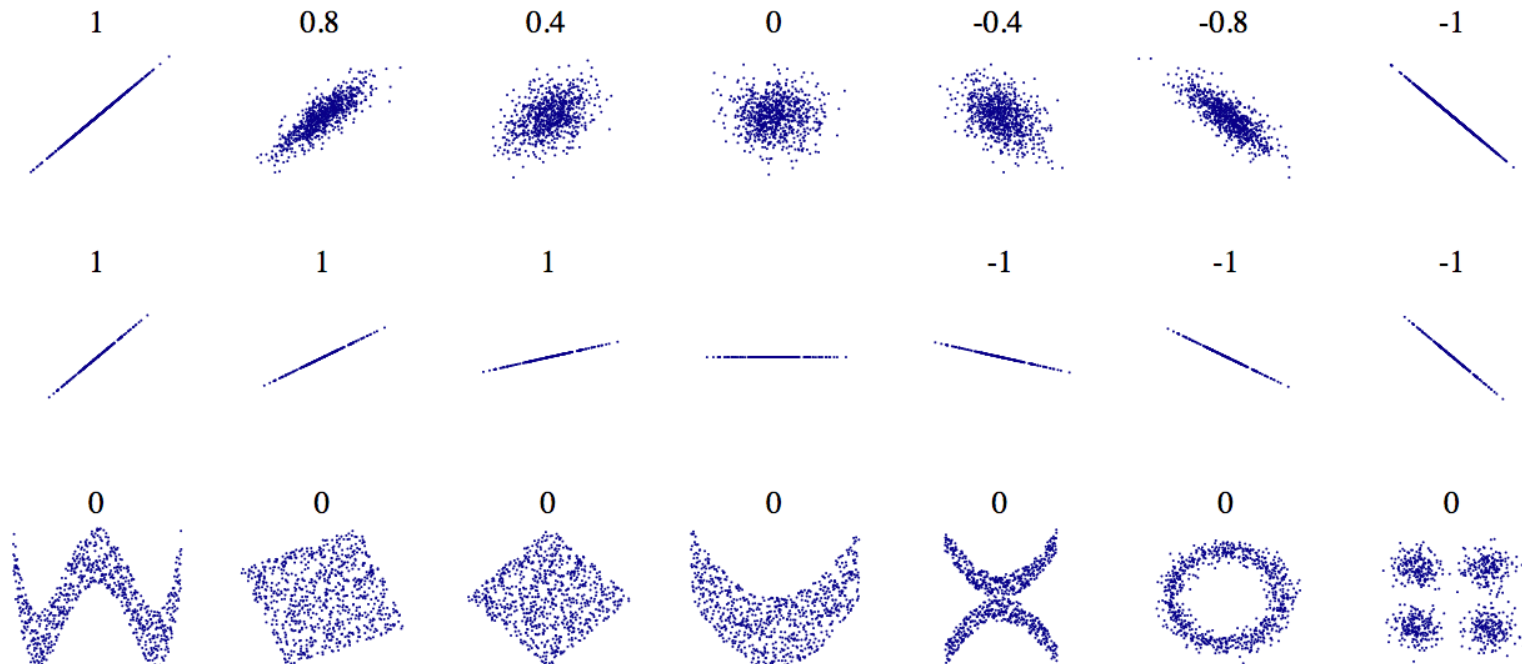
- To visualize the statistics
- E.g.,



❖ Correlations

- Mainly refer to linear correlation (simple for exploration)
- Between a feature and the target, or between two features

❖ Examples ('1' for strong positive correlation)



- ❖ Missing values
 - Blanks, NaNs, or other placeholders
- ❖ Standardization
 - Attributes with varying scales
- ❖ Normalization
- ❖ Imbalanced data
 - Sampling
- ❖ Data duplication
- ❖ Outliers (noise)

- ❖ Discarding the data instances with missing values
 - If the number of missing values is small
- ❖ Discarding the attributes with missing values
 - If the number of missing values is big
- ❖ Imputation
 - Manually: tedious + infeasible?
 - Automatically
 - A global constant
 - **Mean, median, mode**, or other statistics
 - Based on advanced techniques such as **regression, nearest neighbors, matrix factorization**, etc.

Missing Value Imputation



❖ Examples

- Raw data:

$$X = \begin{bmatrix} -1 & 1 \\ 6 & 2 \\ 3 & 3 \\ 3 & 4 \\ 1 & 4 \\ 9 & ? \end{bmatrix}$$

- Mean: 2.8 $\left(\frac{1+2+3+4+4}{5} = 2.8\right)$
- Median: 3 (3 is in the middle)
- Mode: 4 (4 is the most frequent)
- Regression: How? (result: 3.38)

- ❖ Z-score (centered and scaled)

$$x' = \frac{x - \mu}{\sigma}$$

- ❖ Fixed range, e.g., $[0, 1]$ or $[-1, 1]$

- For $[0, 1]$: $x' = \frac{x - x_{min}}{x_{max} - x_{min}}$

- For $[-1, 1]$: ?

- ❖ Non-linear transformation (still order-preserving)

- E.g., Box-Cox transform

- ❖ All transformation should be applied to both training and testing data

Standardization (Cont'd)

❖ Examples

- Raw data:

$$X = \begin{bmatrix} 1 & -1 & 2 \\ 2 & 0 & 0 \\ 0 & 1 & -1 \end{bmatrix}$$

- Z-score:

$$X' = \begin{bmatrix} 1.0 & -1.22 & 1.33 \\ 1.22 & 0.0 & -0.26 \\ -1.22 & 1.22 & -1.06 \end{bmatrix}$$

- Fixed range [0, 1]:

$$X' = \begin{bmatrix} 0.5 & 0.0 & 1.0 \\ 1.0 & 0.5 & 0.33 \\ 0.0 & 1.0 & 0.0 \end{bmatrix}$$

❖ Scaling individual data instances to have unit norm

- Metrics: L_1 norm: Manhattan; L_2 norm: Euclidean
- i.e., the “length” of a feature vector is 1
- The angle information is still preserved

❖ Examples

- Raw data :
$$X = \begin{bmatrix} 1 & -1 & 2 \\ 2 & 0 & 0 \\ 0 & 1 & -1 \end{bmatrix}$$

- L_2 normalized:
$$X' = \begin{bmatrix} 0.40 & -0.40 & 0.81 \\ 1.00 & 0.0 & 0.0 \\ 0.0 & 0.70 & -0.70 \end{bmatrix}$$

- E.g., first row: $\sqrt{0.40^2 + (-0.40)^2 + 0.81^2} \approx 1.0$

- ❖ Sampling: obtaining data instances to represent the whole original data set
- ❖ Undersampling **vs** oversampling
 - The sample size is less than the original size
 - The sample size is greater than the original size
- ❖ With **vs** without replacement
 - Whether an instance is removed from the data set or not
- ❖ Simple random sampling (SRS) **vs** other sampling
 - SRS: each instance has an equal probability of being selected
 - Others: the probability might not be equal, e.g., **stratified sampling** (grouping data first; then sampling from groups)

❖ Data duplication

- Depending on specific ML models, some are sensitive

❖ Outlier (noise)

- Caused by data entry problem, faulty instruments, ...
- The consequence also depends on specific ML models
- E.g., ordinary linear regression is sensitive while tree models are often robust
- Outlier detection itself can be the main ML task, particularly in the cyber security domain

- ❖ Feature generation
 - Manual feature engineering by domain experts
 - Domain experts can identify more informative features
 - Feature interaction and polynomials
- ❖ Feature selection
 - Some features might be redundant or irrelevant
 - Select a subset of (informative) features
- ❖ Feature extraction
 - Transform to another feature space (lower-dimensional)
- ❖ Deep learning
 - Automatic feature representation and extraction



- ❖ To enrich feature representation
 - Particularly for linear models (to capture **non-linearity**)
 - E.g., we used higher orders of polynomial terms in the linear regression running example (still remember?)
- ❖ Feature interaction example


$$\langle X_1, X_2 \rangle \rightarrow \langle 1, X_1, X_2, X_1^2, X_1X_2, X_2^2 \rangle$$

Index	X1	X2
0	0	1
1	2	3

Index	1	X1	X2	X1X1	X1X2	X2X2
0	1	0	1	0	0	1
1	1	2	3	4	6	9

- ❖ Filter approaches: univariate statistics
 - Test if an individual feature has statistically significant relationship with the target
 - E.g., analysis of variance (ANOVA), correlation coefficient
- ❖ Wrapper approaches: feature subset selection
 - Brute-force: try all possible subsets (NOT scalable)
 - Recursive feature elimination
 - Build a model; discard least important feature; repeat
- ❖ Embedded approaches: model-based selection
 - Use ML models to judge feature importance
 - E.g., decision tree models, Lasso regression coefficients

- ❖ Transformative
 - Linear or non-linear projection to another space
 - Dimensionality can usually be reduced
- ❖ Commonly-used techniques
 - E.g., Principal Component Analysis (PCA)

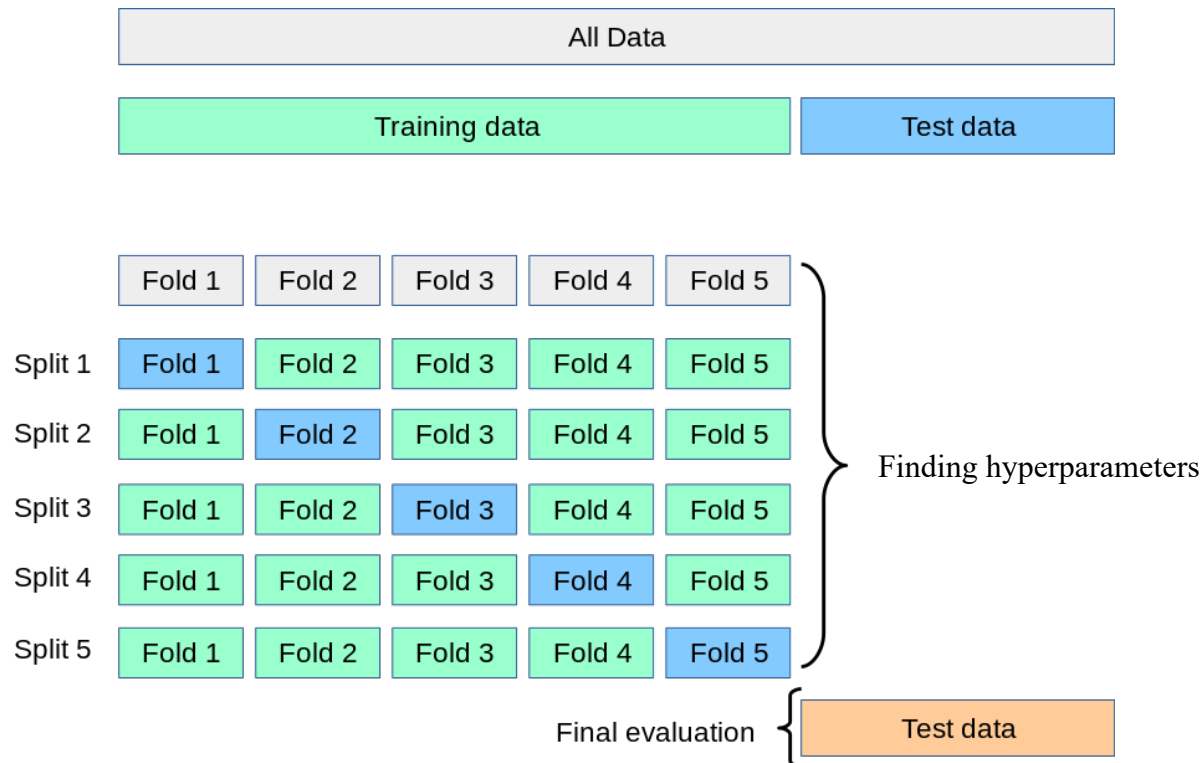


- ❖ Deep learning can automatically perform feature selection and extraction
 - Interpretability might be lost, e.g., feature importance
- ❖ Typical examples
 - Convolutional Neural Network (CNN)
 - Convolution can capture/activate features in images
 - For images/videos
 - Recurrent Neural Network (RNN)
 - Capture temporal/sequential information
 - Long Short-Term Memory networks (LSTM)
 - A special kind of RNN
 - For sequences, audio, time series, and text

- ❖ Determine a machine learning model
 - Supervised vs unsupervised
 - Categorical vs numerical
 - Batch vs online
 - ...
- ❖ Split data into training and testing data
 - Training data for building the model
 - Testing data for validating the learning model
 - Randomly partition the data
 - Can be 80% vs 20%, or 75% vs 25%
 - K-folder cross validation


❖ k -fold cross validation

- Usually, k is 5 or 10
- Extreme case $k = n$, leave one out cross validation



- ❖ Model tuning is a non-trivial job!
 - Model selection
 - Confusing terms
 - Model parameters
 - Model hyperparameters
 - Learning algorithm parameters
 - Tuning is performed w.r.t. **model hyperparameters**
- ❖ Typical tuning methods
 - Manual tuning based on human experts
 - Automated grid search for optimal model hyperparameters
 - Automated random search
 - Often work together with cross validation

- ❖ Deploy the trained model in a production environment
 - Need to save the trained model
 - E.g., using 'joblib' in Scikit-learn library
 - Integrated with other parts of a system
 - E.g., wrap the model within a dedicated web service that your web application can query through a REST API



- Other example is Google Cloud AI Platform

❖ Workflow of Machine Learning Project


- Data Pre-processing and Feature Engineering
- Model Training and Evaluation

❖ Linear Regression Introduction

- Linear Regression Model
- Overfitting and Model Selection

❖ Suppose we have 10 observations

- $\mathcal{D} = \{\langle x_1, t_1 \rangle, \dots, \langle x_{10}, t_{10} \rangle\}$
- Ground truth: $t = \sin(2\pi x)$
- Noise added to t_i



❖ **Regression** task: estimate $t = y(x)$ from D

❖ A function is linear if

$$f(\alpha x_1 + \beta x_2) = \alpha f(x_1) + \beta f(x_2)$$

- α, β are scalars
- Additivity and homogeneity

❖ Linear regression model


$$y(\mathbf{x}, \boldsymbol{\omega}) = \omega_0 + \omega_1 x_1 + \omega_2 x_2 + \cdots + \omega_D x_D$$

- Note: D is the # of features, not the # of instances
- The linearity of $y(\cdot, \boldsymbol{\omega})$ is with respect to $\boldsymbol{\omega}$
- Note: we are learning $\boldsymbol{\omega}$

The Simplest From

- ❖ Dimensionality $M = 1$

$$y(x, \omega) = \omega_0 + \omega_1 x$$



- Can only model a “line”
- ❖ How to handle non-linearity?

❖ Transform \mathbf{x}

$$\mathbf{x} \Rightarrow \mathbf{\Phi}(\mathbf{x}) = (\phi_0(\mathbf{x}), \phi_1(\mathbf{x}), \phi_2(\mathbf{x}), \dots, \phi_M(\mathbf{x}))^T$$

- $\phi_j(x)$, $j \geq 0$ form a family of basic functions
- Transformation/feature function is often non-linear

❖ Linear basis function models

$$y(\mathbf{x}, \boldsymbol{\omega}) = \sum_{j=0}^M \omega_j \phi_j(\mathbf{x}) = \boldsymbol{\omega}^T \mathbf{\Phi}(\mathbf{x})$$

- $\boldsymbol{\omega} = (\omega_0, \dots, \omega_M)^T$, $\mathbf{\Phi} = (\phi_0, \dots, \phi_M)^T$

- ❖ For $0 < j \leq D$, $\phi_j(x) = x_j$; $\phi_0(x) = 1$

$$y(\mathbf{x}, \boldsymbol{\omega}) = \sum_{j=0}^D \omega_j x_j$$

- The original form

- ❖ $\phi_j(x) = x^j$, $0 < j \leq M$

$$y(x, \boldsymbol{\omega}) = \sum_{j=0}^M \omega_j x^j$$

- Power series expansions



❖ How to learn/choose the parameter vector ω ?

- First, define a loss function $L(y(x), t)$
- Then, find out $y(x)$ to minimize the loss (expectation)

$$y^*(\mathbf{x}) = \arg \min_{y(\mathbf{x})} \mathbb{E}[L] = \int \int L(y(\mathbf{x}), t) p(\mathbf{x}, t) d\mathbf{x} dt$$


- E.g., general squared loss function:

$$L(y(\mathbf{x}), t) = \frac{1}{2} \|y(\mathbf{x}) - t\|^2$$

- In practice, we often use the loss function defined on training data

❖ Sum of squared error:

$$E(\omega) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \omega) - t_n\}^2$$



❖ Then, the specific learning problem

$$\omega^* = \arg \min_{\omega} E(\omega)$$

- ❖ Learning problem \rightarrow optimization problem

$$\omega^* = \arg \min_{\omega} E(\omega),$$

$$E(\omega) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \omega) - t_n\}^2$$

- ❖ How to solve? $\frac{\partial E(\omega)}{\partial \omega} = 0$





- Close-form solution can be obtained
- For the basis function model

$$\omega_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

Different Degrees of Polynomials



MACQUARIE
University



- ❖ $M = 0$ or $M = 1$
 - Only passing a couple of points
 - Models are too simple
- ❖ $M = 9$
 - Passing every points but oscillating wildly
 - **Overfitting** problem: Lack of generalisation capability
- ❖ $M = 3$
 - Much better solution
 - Good trade-o has been made

- ❖ Training/empirical error of the trained model \hat{f} on a training dataset with size N , e.g.,

$$E_{emp}(\hat{f}) = \frac{1}{N} \sum_{i=1}^N L(y_i, \hat{f}(x_i))$$

- Learning process usually tries to minimize this error

- ❖ Test error on a test data set with size N' , e.g.,

$$E_{test}(\hat{f}) = \frac{1}{N'} \sum_{i=1}^{N'} L(y_i, \hat{f}(x_i))$$


- Test error indicates the **generalisation capability** of a learned model

Overfitting Issue

❖ Test errors with respect to M

- Test data containing 100 points
- Root-Mean-Square (RMS) error

$$E_{RMS} = \sqrt{2E(\omega^*)/N}$$



- ❖ Observations on M
 - Different models generated by varying M
 - M controls the learning capability of the model
 - Unlike model parameter ω , M can NOT be learned from data (**hyperparameter**)
- ❖ Model selection
 - Models generated by different learning algorithms
 - Models generated by different algorithm parameter (model hyperparameter) configurations
- ❖ Model selection is a non-trivial job!

❖ Another way to the overfitting issue

- Penalizing large coefficients in error function
- E.g., quadratic regularizer




$$\tilde{E}(\omega) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \omega) - t_n\}^2 + \frac{\lambda}{2} \|\omega\|^2,$$

$$\|\omega\|^2 = \omega^T \omega = \omega_0^2 + \omega_1^2 + \cdots + \omega_M^2$$

- λ governs the trade-offs between the two terms
- A.k.a. ridge regression
- Close-form solution can be obtained
- Lasso regression: the regularizer is $\|\omega\|_{L_1}$

Effects of λ

❖ Performance of different λ when $M = 9$





- ❖ Data preparation
- ❖ Training/test data split
 - `x_train, y_train, x_test, y_test`
- ❖ Train the model
 - Build a model: `model = xxxModel.xxxModel()`
 - Train a model: `model.fit(x_train, y_train)`
 - Validate: `cross_validate(xxxmodel, x_train, y_train, k, scoring)`
- ❖ Test the learning model
 - Test the model: `y_pred = xxxModel.predict(x_test)`
 - Can use `y_pred` and `y_test` to calculate the error, e.g., `mean_squared_error(y_test, y_pred)`